





MoE variants vs baselines (7.86B tokens, same architecture family)



† Boltzmann rows: routing weights $p_i = \text{softmax}(E_i/\tau)$ derived from energy scores (no learned router).
Dashed line connects h1_topk_egpt_moe ↔ h1_boltz_fullsize (same H-series architecture, different routing).