


Consecutive-layer cosine similarity — d=1024, 24×1 (V9 GPT, V19 EGrad, V20 EDesc)
cos_sim(layer[i], layer[i+1]) | higher = adjacent layers produce more similar outputs

attn_out (pre-residual attention)



ffwd_out (pre-residual FFN)

