


Consecutive-layer cosine similarity — d=768, 12×1 (V0 GPT, V1 EGPT, V15 EGrad, V16 EDesc)
cos_sim(layer[i], layer[i+1]) | higher = adjacent layers produce more similar outputs

attn_out (pre-residual attention)



ffwd_out (pre-residual FFN)

