

Inter-layer cosine similarity — d=1024, 24×1 (V9 GPT, V19 EGrad, V20 EDesc)
attn_out = pre-residual attention output; ffwd_out = pre-residual FFN output

