

Inter-layer cosine similarity — d=768, 12×1 (V0 GPT, V1 EGPT, V15 EGrad, V16 EDesc)
attn_out = pre-residual attention output; ffwd_out = pre-residual FFN output

