

Accuracy vs Compute: Mixed heads are strictly more efficient than pure EGPT at same architecture family

