


Per-task accuracy: Mixed vs pure architectures

12×1 family (~143–162M)



6×2 recurrent family (~110–120M)

